

1 Phase 4 Project - Image Classification

In this project, we will aim to use artificial neural networks to identify pneumonia by looking at x-ray images of patients' lungs.

1.1 Obtain

The data for this project comes from Kaggle. <https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia> (<https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>)

We will download the images into the './data' folder, then use Keras' ImageDataGenerator to feed the images into our models.

```
In [1]: import tensorflow as tf
        from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

executed in 5.66s, finished 11:33:34 2021-04-30

```
In [2]: # Instantiate ImageDataGenerators for train, validation, and test data
        train_data = ImageDataGenerator(rescale=1/255,
                                         shear_range=0.2,
                                         zoom_range=0.2,
                                         horizontal_flip=True)

        val_data = ImageDataGenerator(rescale=1/255)
        test_data = ImageDataGenerator(rescale=1/255)

        # Feed the data
        train_generator = train_data.flow_from_directory('./data/train',
                                                         color_mode='rgb',
                                                         target_size=(224,224),
                                                         batch_size=32,
                                                         class_mode='binary')

        val_generator = val_data.flow_from_directory('./data/val',
                                                      color_mode='rgb',
                                                      target_size=(224,224),
                                                      batch_size=16,
                                                      shuffle=False,
                                                      class_mode='binary')

        test_generator = test_data.flow_from_directory('./data/test',
                                                        color_mode='rgb',
                                                        target_size=(224,224),
                                                        batch_size=32,
                                                        shuffle=False,
                                                        class_mode='binary')
```

executed in 732ms, finished 11:33:34 2021-04-30

Found 5216 images belonging to 2 classes.
 Found 16 images belonging to 2 classes.
 Found 624 images belonging to 2 classes.





1.2 Scrub

There isn't much to be done for this dataset in terms of scrubbing. Each image is saved into folders based on what set it is part of (train, val, or test) and what its label is (normal or pneumonia). To go through the images and look for inconsistencies would require medical knowledge beyond my training.



1.3 Examine

```
In [3]: import matplotlib.pyplot as plt  
        %matplotlib inline
```

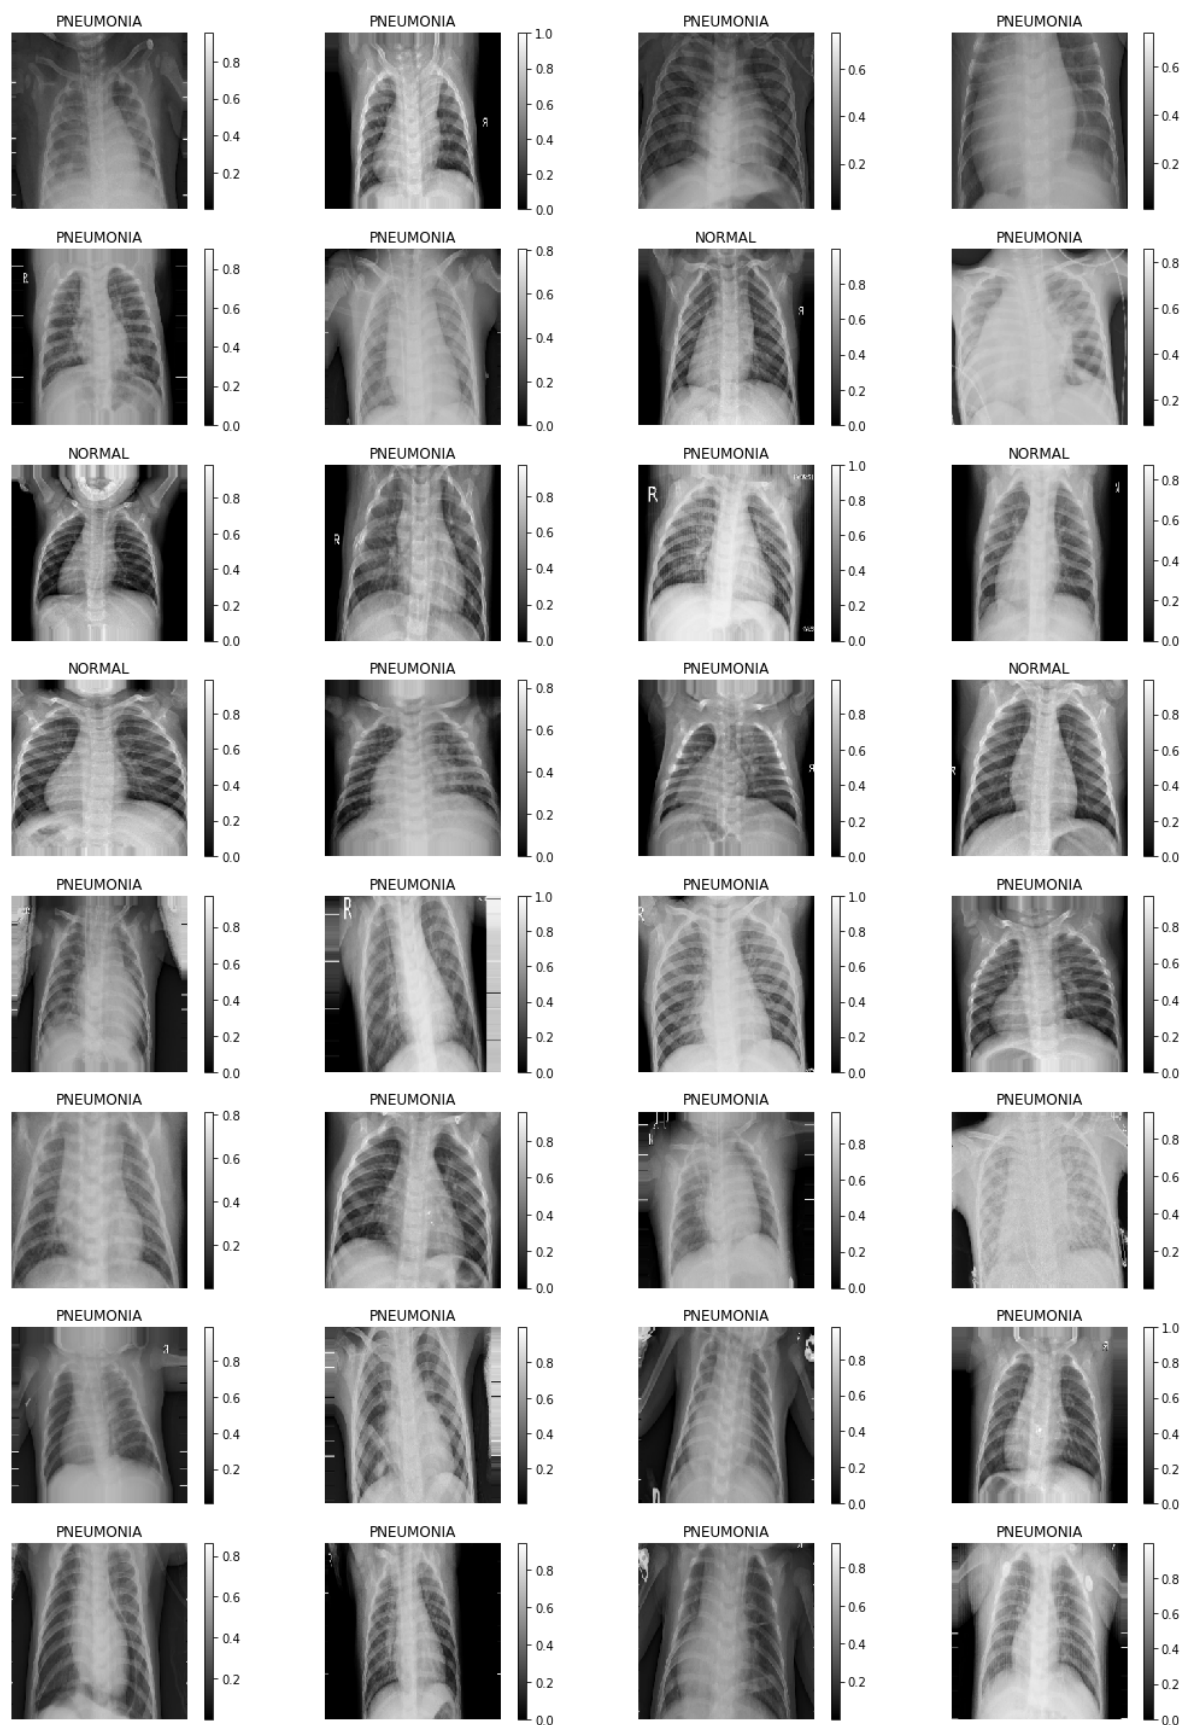
executed in 280ms, finished 11:33:35 2021-04-30

We'll use our train_generator to load up some images for us to look at.

```
In [4]: img, labels = train_generator.next()

plt.figure(figsize=(15,20))
▼ for i in range(len(img)):
    label = 'NORMAL'
▼ if labels[i] == 1:
    label = 'PNEUMONIA'
    plt.subplot(8,4,i+1)
    plt.title(label)
    plt.imshow(img[i], cmap='gray')
    plt.axis('off')
    plt.colorbar()
plt.tight_layout()
```

executed in 3.62s, finished 11:33:38 2021-04-30



Just at a first glance, the 'pneumonia' images appear to be a little fuzzier. The 'normal' images appear to have a more defined black area between the ribs. However, it's not something I would be comfortable trying to differentiate on my own, so let's try out some neural networks!

1.4 Model

```
In [5]: from tensorflow.keras.models import *
        from tensorflow.keras.layers import *
        from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping

        from sklearn.metrics import confusion_matrix, classification_report
```

executed in 745ms, finished 11:33:39 2021-04-30

```
In [6]: # Adding some callbacks to avoid training for too long
        def init_callbacks():
            early_stopping = EarlyStopping(monitor='val_loss', patience=5)
            model_checkpoint = ModelCheckpoint('best_model.h5', monitor='val_loss',
                                              verbose=1, save_best_only=True,
                                              save_weights_only=True)

            callbacks = [early_stopping, model_checkpoint]
            return callbacks
```

executed in 15ms, finished 11:33:39 2021-04-30

1.4.1 Model 1

Let's start with a model with a single hidden layer. We'll use 128 nodes with 'relu' activation.

```
In [7]: model_1 = Sequential()
        model_1.add(Flatten(input_shape=(224,224,3)))
        model_1.add(Dense(128, activation='relu'))
        model_1.add(Dense(1, activation='sigmoid'))

        model_1.compile(loss='binary_crossentropy',
                        optimizer='adam',
                        metrics=['accuracy'])
        model_1.summary()
```

executed in 286ms, finished 11:33:39 2021-04-30

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 150528)	0
dense (Dense)	(None, 128)	19267712
dense_1 (Dense)	(None, 1)	129
=====		
Total params: 19,267,841		
Trainable params: 19,267,841		
Non-trainable params: 0		

```
In [8]: history_1 = model_1.fit(train_generator,
                                epochs=25,
                                validation_data=val_generator,
                                callbacks=init_callbacks())
```

executed in 21m 24s, finished 11:55:04 2021-04-30

Epoch 1/25

163/163 [=====] - 85s 518ms/step - loss: 10.2409 - accuracy: 0.7107 - val_loss: 1.1253 - val_accuracy: 0.6250

Epoch 00001: val_loss improved from inf to 1.12532, saving model to best_model.h5

Epoch 2/25

163/163 [=====] - 83s 510ms/step - loss: 0.5110 - accuracy: 0.8483 - val_loss: 1.3954 - val_accuracy: 0.6250

Epoch 00002: val_loss did not improve from 1.12532

Epoch 3/25

163/163 [=====] - 79s 484ms/step - loss: 0.7264 - accuracy: 0.8335 - val_loss: 0.4280 - val_accuracy: 0.8125

Epoch 00003: val_loss improved from 1.12532 to 0.42795, saving model to best_model.h5

Epoch 4/25

163/163 [=====] - 79s 483ms/step - loss: 0.5145 - accuracy: 0.8560 - val_loss: 0.5614 - val_accuracy: 0.8125

Epoch 00004: val_loss did not improve from 0.42795

Epoch 5/25

163/163 [=====] - 79s 484ms/step - loss: 0.4260 - accuracy: 0.8654 - val_loss: 0.4773 - val_accuracy: 0.8750

Epoch 00005: val_loss did not improve from 0.42795

Epoch 6/25

163/163 [=====] - 80s 488ms/step - loss: 0.3632 - accuracy: 0.8713 - val_loss: 0.4110 - val_accuracy: 0.8125

Epoch 00006: val_loss improved from 0.42795 to 0.41102, saving model to best_model.h5

Epoch 7/25

163/163 [=====] - 80s 493ms/step - loss: 0.5262 - accuracy: 0.8736 - val_loss: 0.5517 - val_accuracy: 0.8125

Epoch 00007: val_loss did not improve from 0.41102

Epoch 8/25

163/163 [=====] - 81s 496ms/step - loss: 0.2719 - accuracy: 0.8879 - val_loss: 1.7139 - val_accuracy: 0.5625

Epoch 00008: val_loss did not improve from 0.41102

Epoch 9/25

163/163 [=====] - 81s 494ms/step - loss: 0.4088 - accuracy: 0.8603 - val_loss: 0.4153 - val_accuracy: 0.8125

Epoch 00009: val_loss did not improve from 0.41102

Epoch 10/25

163/163 [=====] - 80s 491ms/step - loss: 0.3270 - accuracy: 0.8879 - val_loss: 0.41102 - val_accuracy: 0.8125

racy: 0.8788 - val_loss: 0.4605 - val_accuracy: 0.8125

Epoch 00010: val_loss did not improve from 0.41102

Epoch 11/25

163/163 [=====] - 80s 492ms/step - loss: 0.2860 - accuracy: 0.9011 - val_loss: 0.3609 - val_accuracy: 0.8750

Epoch 00011: val_loss improved from 0.41102 to 0.36087, saving model to best_model.h5

Epoch 12/25

163/163 [=====] - 79s 485ms/step - loss: 0.2330 - accuracy: 0.9082 - val_loss: 0.4875 - val_accuracy: 0.8125

Epoch 00012: val_loss did not improve from 0.36087

Epoch 13/25

163/163 [=====] - 80s 492ms/step - loss: 0.2534 - accuracy: 0.9000 - val_loss: 0.4168 - val_accuracy: 0.7500

Epoch 00013: val_loss did not improve from 0.36087

Epoch 14/25

163/163 [=====] - 79s 487ms/step - loss: 0.2462 - accuracy: 0.8994 - val_loss: 0.5039 - val_accuracy: 0.8125

Epoch 00014: val_loss did not improve from 0.36087

Epoch 15/25

163/163 [=====] - 80s 488ms/step - loss: 0.2433 - accuracy: 0.8950 - val_loss: 0.5235 - val_accuracy: 0.8125

Epoch 00015: val_loss did not improve from 0.36087

Epoch 16/25

163/163 [=====] - 79s 486ms/step - loss: 0.2679 - accuracy: 0.8914 - val_loss: 0.4259 - val_accuracy: 0.6875

Epoch 00016: val_loss did not improve from 0.36087

Because of our callbacks, the model stopped training early because the loss was not improving on the validation set.

Let's make a quick function to plot out the model's performance across epochs.


```
In [9]: def plot_model(hist):
        '''
        input: fitted model
        output: plots of accuracy and loss
        '''

        x = range(1,len(hist.history['loss'])+1)
        fig,ax = plt.subplots(2,1,figsize=(12,8))
        ax[0].plot(x, hist.history['accuracy'], label='acc')
        ax[0].plot(x, hist.history['val_accuracy'], label='val acc')
        ax[0].legend()
        ax[0].set_title('Accuracy')
        ax[1].plot(x, hist.history['loss'], label='loss')
        ax[1].plot(x, hist.history['val_loss'], label='val loss')
        ax[1].legend()
        ax[1].set_title('Loss')
        plt.show()
```

executed in 13ms, finished 11:55:04 2021-04-30

In [10]: plot_model(history_1)

executed in 266ms, finished 11:55:04 2021-04-30



Those lines really started diverging quickly. Let's load up the best weights and check our performance on the test set.


```
In [11]: model_1.save_weights('last_model.h5')
         model_1.evaluate(test_generator)
```

executed in 5.03s, finished 11:55:09 2021-04-30

20/20 [=====] - 4s 205ms/step - loss: 0.3286 - accurac
y: 0.8526

Out[11]: [0.32862070202827454, 0.8525640964508057]

```
In [12]: model_1.load_weights('best_model.h5')
         model_1.evaluate(test_generator)
```

executed in 4.95s, finished 11:55:14 2021-04-30

20/20 [=====] - 4s 202ms/step - loss: 0.3138 - accurac
y: 0.8590

Out[12]: [0.3138106167316437, 0.8589743375778198]

```
In [13]: predictions = model_1.predict(test_generator)
         y_pred = (predictions > 0.5).astype('int')

         y_true = test_generator.classes
```

executed in 4.91s, finished 11:55:19 2021-04-30

```
In [14]: print(classification_report(y_true, y_pred))
         print(confusion_matrix(y_true, y_pred))
```

executed in 14ms, finished 11:55:19 2021-04-30

	precision	recall	f1-score	support
0	0.84	0.77	0.80	234
1	0.87	0.91	0.89	390
accuracy			0.86	624
macro avg	0.85	0.84	0.85	624
weighted avg	0.86	0.86	0.86	624

```
[[181  53]
 [ 35 355]]
```

Our base model gets us 86% accuracy with a 91% recall on positive cases. This is a good start. Let's see if we can improve our model.

▼ 1.4.2 Model 2 - Deeper

For the next model, we'll try going a little deeper with three hidden layers.

```
In [15]: model_2 = Sequential()
model_2.add(Flatten(input_shape=(224,224,3)))
model_2.add(Dense(64, activation='relu'))
model_2.add(Dense(64, activation='relu'))
model_2.add(Dense(64, activation='relu'))
model_2.add(Dense(1, activation='sigmoid'))

▼ model_2.compile(loss='binary_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])
model_2.summary()
```

executed in 109ms, finished 11:55:19 2021-04-30

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
flatten_1 (Flatten)	(None, 150528)	0
dense_2 (Dense)	(None, 64)	9633856
dense_3 (Dense)	(None, 64)	4160
dense_4 (Dense)	(None, 64)	4160
dense_5 (Dense)	(None, 1)	65
=====		
Total params: 9,642,241		
Trainable params: 9,642,241		
Non-trainable params: 0		

```
In [16]: history_2 = model_2.fit(train_generator,  
                                epochs=25,  
                                validation_data=val_generator,  
                                callbacks=init_callbacks())
```

executed in 8m 58s, finished 12:04:18 2021-04-30

Epoch 1/25

163/163 [=====] - 76s 466ms/step - loss: 3.8824 - accuracy: 0.6885 - val_loss: 0.9281 - val_accuracy: 0.6250

Epoch 00001: val_loss improved from inf to 0.92809, saving model to best_model.h5

Epoch 2/25

163/163 [=====] - 76s 468ms/step - loss: 0.3942 - accuracy: 0.8427 - val_loss: 0.4331 - val_accuracy: 0.8750

Epoch 00002: val_loss improved from 0.92809 to 0.43315, saving model to best_model.h5

Epoch 3/25

163/163 [=====] - 77s 470ms/step - loss: 0.3546 - accuracy: 0.8508 - val_loss: 0.4934 - val_accuracy: 0.8125

Epoch 00003: val_loss did not improve from 0.43315

Epoch 4/25

163/163 [=====] - 77s 472ms/step - loss: 0.3205 - accuracy: 0.8655 - val_loss: 0.7093 - val_accuracy: 0.6250

Epoch 00004: val_loss did not improve from 0.43315

Epoch 5/25

163/163 [=====] - 77s 470ms/step - loss: 0.2888 - accuracy: 0.8773 - val_loss: 0.6411 - val_accuracy: 0.6875

Epoch 00005: val_loss did not improve from 0.43315

Epoch 6/25

163/163 [=====] - 78s 476ms/step - loss: 0.3037 - accuracy: 0.8755 - val_loss: 0.5380 - val_accuracy: 0.7500

Epoch 00006: val_loss did not improve from 0.43315


Epoch 7/25

163/163 [=====] - 77s 475ms/step - loss: 0.2961 - accuracy: 0.8761 - val_loss: 0.4452 - val_accuracy: 0.7500

Epoch 00007: val_loss did not improve from 0.43315

In [17]: `plot_model(history_2)`

executed in 253ms, finished 12:04:18 2021-04-30



In [18]: `model_2.save_weights('last_model.h5')`
`model_2.evaluate(test_generator)`

executed in 4.92s, finished 12:04:23 2021-04-30

20/20 [=====] - 4s 201ms/step - loss: 0.3249 - accuracy: 0.8542

Out[18]: [0.32494595646858215, 0.8541666865348816]

In [19]:

```
model_2.load_weights('best_model.h5')
model_2.evaluate(test_generator)
```

executed in 4.86s, finished 12:04:28 2021-04-30

20/20 [=====] - 4s 200ms/step - loss: 0.3420 - accuracy: 0.8413

Out[19]: [0.3419956564903259, 0.8413461446762085]

In [20]:

```
y_pred = model_2.predict(test_generator)
y_pred = (y_pred > 0.5).astype('int32')

print(classification_report(y_true, y_pred))
print(confusion_matrix(y_true, y_pred))
```

executed in 4.76s, finished 12:04:32 2021-04-30

	precision	recall	f1-score	support
0	0.82	0.74	0.78	234
1	0.85	0.90	0.88	390
accuracy			0.84	624
macro avg	0.84	0.82	0.83	624
weighted avg	0.84	0.84	0.84	624

```
[[174  60]
 [ 39 351]]
```

▼ 1.4.3 Model 3 - Convolutional Neural Network

Convolutional neural networks have a reputation for being effective at image classification, so let's see if we can use the power of CNNs to improve our performance.

```
In [21]: model_3 = Sequential()
model_3.add(Conv2D(128, kernel_size=3, activation='relu',
                  input_shape=(224,224,3)))
model_3.add(Conv2D(64, kernel_size=3, activation='relu'))
model_3.add(Conv2D(32, kernel_size=3, activation='relu'))
model_3.add(Flatten())
model_3.add(Dense(1, activation='sigmoid'))

model_3.compile(loss='binary_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])
model_3.summary()
```

executed in 91ms, finished 12:04:32 2021-04-30

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 222, 222, 128)	3584
conv2d_1 (Conv2D)	(None, 220, 220, 64)	73792
conv2d_2 (Conv2D)	(None, 218, 218, 32)	18464
flatten_2 (Flatten)	(None, 1520768)	0
dense_6 (Dense)	(None, 1)	1520769
Total params: 1,616,609		
Trainable params: 1,616,609		
Non-trainable params: 0		

Convolutional Neural Networks take much more time to train, so I'm going to train this one in Google Colab and then load the weights after it's completed.

```
In [22]: model_3 = tf.keras.models.load_model('cnn/content/cnn')
```

executed in 379ms, finished 12:04:33 2021-04-30

```
In [23]: model_3.evaluate(test_generator)
```

executed in 19.2s, finished 12:04:52 2021-04-30

20/20 [=====] - 19s 907ms/step - loss: 0.5595 - accuracy: 0.7804

```
Out[23]: [0.5594589710235596, 0.7804487347602844]
```

```
In [24]: y_pred = model_3.predict(test_generator)
y_pred = (y_pred > 0.5).astype('int')

print(classification_report(y_true, y_pred))
print(confusion_matrix(y_true, y_pred, normalize='true'))
```

executed in 19.9s, finished 12:05:12 2021-04-30

	precision	recall	f1-score	support
0	0.90	0.47	0.61	234
1	0.75	0.97	0.85	390
accuracy			0.78	624
macro avg	0.83	0.72	0.73	624
weighted avg	0.81	0.78	0.76	624


```
[[0.46581197 0.53418803]
 [0.03076923 0.96923077]]
```

▼ 1.4.4 Model 4 - More Complex CNN

This model will add MaxPooling2D layers for downsampling and Dropout layers to help reduce overfitting.

```
In [25]: model_4 = Sequential()
model_4.add(Conv2D(32, kernel_size=(3,3), activation='relu',
                  input_shape=(224,224,3)))
model_4.add(Conv2D(64, kernel_size=(3,3), activation='relu'))
model_4.add(MaxPooling2D(pool_size=(2,2)))
model_4.add(Dropout(0.25))
model_4.add(Conv2D(128, kernel_size=(3,3), activation='relu'))
model_4.add(MaxPooling2D())
model_4.add(Dropout(0.25))
model_4.add(Flatten())
model_4.add(Dense(64, activation='relu'))
model_4.add(Dropout(0.5))
model_4.add(Dense(1, activation='sigmoid'))

model_4.compile(loss='binary_crossentropy',
                optimizer='adam',
                metrics=['accuracy'])

model_4.summary()
```

executed in 203ms, finished 12:05:12 2021-04-30

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 222, 222, 32)	896
conv2d_4 (Conv2D)	(None, 220, 220, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 110, 110, 64)	0
dropout (Dropout)	(None, 110, 110, 64)	0
conv2d_5 (Conv2D)	(None, 108, 108, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 128)	0
dropout_1 (Dropout)	(None, 54, 54, 128)	0
flatten_3 (Flatten)	(None, 373248)	0
dense_7 (Dense)	(None, 64)	23887936
dropout_2 (Dropout)	(None, 64)	0
dense_8 (Dense)	(None, 1)	65
Total params: 23,981,249		
Trainable params: 23,981,249		
Non-trainable params: 0		

With 23.9 million trainable parameters, we're going to train this one in Colab again.

In [26]: `model_4 = tf.keras.models.load_model('cnn2/content/cnn2')`

executed in 1.43s, finished 12:05:14 2021-04-30

In [27]: `model_4.evaluate(test_generator)`

executed in 10.9s, finished 12:05:24 2021-04-30

20/20 [=====] - 10s 491ms/step - loss: 0.5111 - accuracy: 0.8157

Out[27]: [0.5110576152801514, 0.8157051205635071]

In [28]: `y_pred = model_4.predict(test_generator)`
`y_pred = (y_pred > 0.5).astype('int32')`
`print(classification_report(y_true, y_pred))`
`print(confusion_matrix(y_true, y_pred, normalize='true'))`

executed in 10.7s, finished 12:05:35 2021-04-30

	precision	recall	f1-score	support
0	0.98	0.52	0.68	234
1	0.78	0.99	0.87	390
accuracy			0.82	624
macro avg	0.88	0.76	0.78	624
weighted avg	0.85	0.82	0.80	624

```
[[0.52136752 0.47863248]
 [0.00769231 0.99230769]]
```

This model achieves 99% recall with 82% accuracy. It's not our most accurate model, but I think the boost in recall makes up for it. With our model identifying 99% of patients with pneumonia, a few false positives is an acceptable compromise.

▼ 1.5 iNterpret

Let's quickly compare our models.

In [29]: `def get_preds(model):`
 `'''returns a models predictions for the test set'''`
 `y_pred = model.predict(test_generator)`
 `y_pred = (y_pred > 0.5).astype('int')`
 `return y_pred`

executed in 14ms, finished 12:05:35 2021-04-30

```
In [30]: models = [model_1, model_2, model_3, model_4]
        for i in range(4):
            print(f'Model {i+1}:')
            y_pred = get_preds(models[i])
            print(classification_report(y_true, y_pred))
            print()
            print(confusion_matrix(y_true, y_pred))
            print()
```

executed in 39.9s, finished 12:06:15 2021-04-30

Model 1:

	precision	recall	f1-score	support
0	0.84	0.77	0.80	234
1	0.87	0.91	0.89	390
accuracy			0.86	624
macro avg	0.85	0.84	0.85	624
weighted avg	0.86	0.86	0.86	624

```
[[181 53]
 [ 35 355]]
```

Model 2:

	precision	recall	f1-score	support
0	0.82	0.74	0.78	234
1	0.85	0.90	0.88	390
accuracy			0.84	624
macro avg	0.84	0.82	0.83	624
weighted avg	0.84	0.84	0.84	624

```
[[174 60]
 [ 39 351]]
```

Model 3:

	precision	recall	f1-score	support
0	0.90	0.47	0.61	234
1	0.75	0.97	0.85	390
accuracy			0.78	624
macro avg	0.83	0.72	0.73	624
weighted avg	0.81	0.78	0.76	624

```
[[109 125]
 [ 12 378]]
```

Model 4:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.98	0.52	0.68	234
1	0.78	0.99	0.87	390
accuracy			0.82	624
macro avg	0.88	0.76	0.78	624
weighted avg	0.85	0.82	0.80	624

```
[[122 112]
 [ 3 387]]
```

With class support of 234 negative to 390 positive, we would be getting ~50% accuracy with random guesses and ~60% accuracy with all-majority-class predictions. So every model has learned something and shows improved performance over random guessing.

With the problem at hand, we want a model that shows high recall on positive cases (predictions identify positive cases) and high precision on negative cases (negative predictions are true). We got a big jump on those metrics when we moved to the convolutional neural network, although the overall accuracy did take a bit of a hit. However, with our more complex CNN, we were able to improve accuracy to 82% with 99% positive recall and 98% negative precision.

2 Conclusion

We were able to get good results using Convolutional Neural Networks. While the accuracy was a bit lower than that of the more traditional neural networks, it did improve the positive recall significantly.

2.1 Recommendations

If pursuing this topic to further improve results, I would make the following recommendations.

1. More Data - The dataset used here is only a portion of a larger dataset. While the larger dataset does not necessarily guarantee better performance, more data will help to reduce overfitting because of increased variance in the training data. A larger dataset should also give a statistically stronger representation of real-world implementation.
2. Pretrained Models - There are a number of pretrained models out there which may be effective for this problem. Many of them are included with Keras. We did not explore any of those models in this notebook, but some research into those models may find some that improve performance without the work of constructing your own models.
3. Address Class Imbalance - When constructing these models, I did not address the issue of class imbalance. While I think the results are satisfactory, addressing the issue of class imbalance early (perhaps by constructing a custom loss function) might improve performance.

2.2 Future Work

Given more time with this data, I'd like to explore some of the popular CNN architectures. Whether or not they improve performance, I would like to get a look into the structure of these models to get a better idea of how convolutional neural networks accomplish their tasks.

In []: